

STUDENT 26: INSURANCE CLAIM & POLICY

=====

A1 :Fragment & Recombine Main Fact (≤ 10 rows)

WHAT TO DO

1. Create horizontally fragmented tables Claim_A on Node_A and Claim_B on Node_B using a deterministic rule (HASH or RANGE on a natural key).
2. Insert a TOTAL of ≤ 10 committed rows split across the two fragments (e.g., 5 on Node_A and 5 on Node_B). Reuse these rows for all remaining tasks.
3. On Node_A, create view Claim_ALL as UNION ALL of Claim_A and Claim_B@proj_link.
4. Validate with COUNT(*) and a checksum on a key column (e.g., SUM(MOD(primary_key,97))) :results must match fragments vs Claim_ALL.

EXPECTED OUTPUT

- ✓ - DDL for Claim_A and Claim_B; population scripts with ≤ 10 total committed rows.
- ✓ - CREATE DATABASE LINK proj_link ... (shown).
- ✓ - CREATE VIEW Claim_ALL ... UNION ALL ... (shown).
- ✓ - Matching COUNT(*) and checksum between fragments vs Claim_ALL (evidence screenshot).

A1.1a

```
001
002 ---CREATING FRAGMENTED TABLE (Claim_A)
003
004 CREATE TABLE IF NOT EXISTS Claim_A (
005     ClaimID SERIAL PRIMARY KEY,      -- this Unique identifier for each claim
006     PolicyID INT NOT NULL,           -- Foreign key referencing the related policy
007     Type VARCHAR(50),                -- Type of claim (e.g., Accident, Theft, Fire)
008     Status VARCHAR(20),              -- Current status of the claim (e.g., Pending, Approved,
009     ClaimedAmount DECIMAL(10,2)      -- Amount claimed by the policyholder
010 );
011
012 ---INSERTING 5 ROW (PolicyID ≤ 3)
013
014 INSERT INTO Claim_A (PolicyID, Type, Status, ClaimedAmount) VALUES
015 (1, 'Health', 'Approved', 80000), -- CLAIMS FOR HEALTH POLICY THAT WAS A. PROVED
016 (2, 'Auto', 'Pending', 60000),    -- claims for auto policy that is pending
017 (3, 'Home', 'Closed', 90000).    -- Home insurance claims that was closed
```

Data Output [Messages](#) [Notifications](#)

NOTICE: relation "claim_a" already exists, skipping
CREATE TABLE

Query returned successfully in 10 secs 995 msec.

A1.1b

```
223
224
225 ---CREATING FRAGMENTED TABLE (Claim_B)
226 CREATE TABLE IF NOT EXISTS Claim_B (
227     ClaimID SERIAL PRIMARY KEY,
228     PolicyID INT NOT NULL,
229     Type VARCHAR(50),
230     Status VARCHAR(20),
231     ClaimedAmount DECIMAL(10,2)
232 );
233
234 ---INSERTING 5 ROW (PolicyID ≤ 3)
235
236 INSERT INTO Claim_B (PolicyID, Type, Status, ClaimedAmount) VALUES
237 (4, 'Health', 'Approved', 50000),
238 (5, 'Home', 'Pending', 70000),
239 (6, 'Life', 'Approved', 110000),
240 (7, 'Auto', 'Closed', 45000),
241 (8, 'Travel', 'Approved', 60000);
242 SELECT * FROM Claim B ;
```

Data Output [Messages](#) [Notifications](#)

INSERT 0 5

Query returned successfully in 711 msec.

A1.2a

```

209
210 ---INSERTING 5 ROW (PolicyID ≤ 3)
211
212 INSERT INTO Claim_A (PolicyID, Type, Status, ClaimedAmount) VALUES
213 (1, 'Health', 'Approved', 80000), -- CLAIMS FOR HEALTH POLICY THAT WAS A. PROVED
214 (2, 'Auto', 'Pending', 60000),    -- claims for auto policy that is pending
215 (3, 'Home', 'Closed', 90000),    |-- Home insurance claims that was closed
216 (1, 'Life', 'Approved', 120000),  --- Life insurance claim approved for payout
217 (2, 'Travel', 'Approved', 40000); --Travel insurance claim successfully approved
218
219 SELECT * FROM Claim_A ;
220

```

Data Output Messages Notifications

| | claimid [PK] integer | policyid integer | type character varying (50) | status character varying (20) | claimedamount numeric (10,2) |
|---|-------------------------|---------------------|--------------------------------|----------------------------------|---------------------------------|
| 1 | 1 | 1 | Health | Approved | 80000.00 |
| 2 | 2 | 2 | Auto | Pending | 60000.00 |
| 3 | 3 | 3 | Home | Closed | 90000.00 |
| 4 | 4 | 1 | Life | Approved | 120000.00 |
| 5 | 5 | 2 | Travel | Approved | 40000.00 |

A1.2b

```

230
231 ---INSERTING 5 ROW (PolicyID ≤ 3)
232
233 INSERT INTO Claim_B (PolicyID, Type, Status, ClaimedAmount) VALUES
234 (4, 'Health', 'Approved', 50000),
235 (5, 'Home', 'Pending', 70000),
236 (6, 'Life', 'Approved', 110000),
237 (7, 'Auto', 'Closed', 45000),
238 (8, 'Travel', 'Approved', 60000);
239
240 SELECT * FROM Claim_B ;
241
242
243 -- Enable the dblink extension (only needs to be done once per database)
244 -- Establish a connection to the remote database (BranchDB_B)
245 -- 'proj_link' is a connection name you can reuse later in queries
246 -- Replace connection info with your actual credentials if needed
247
248 CREATE EXTENSION IF NOT EXISTS dblink;
249 SELECT dblink_connect(
250     'proj_link',

```

Data Output Messages Notifications

| | claimid [PK] integer | policyid integer | type character varying (50) | status character varying (20) | claimedamount numeric (10,2) |
|---|-------------------------|---------------------|--------------------------------|----------------------------------|---------------------------------|
| 1 | 1 | 4 | Health | Approved | 50000.00 |
| 2 | 2 | 5 | Home | Pending | 70000.00 |
| 3 | 3 | 6 | Life | Approved | 110000.00 |
| 4 | 4 | 7 | Auto | Closed | 45000.00 |
| 5 | 5 | 8 | Travel | Approved | 60000.00 |

A1.3

```

248 --Create the Combined View (Claim_ALL)
249
250 CREATE OR REPLACE VIEW Claim_ALL AS
251 SELECT ClaimID, PolicyID, Type, Status, ClaimedAmount
252 FROM Claim_A
253
254 UNION ALL
255
256 SELECT ClaimID, PolicyID, Type, Status, ClaimedAmount
257 FROM dblink('proj_link', 'SELECT ClaimID, PolicyID, Type, Status, ClaimedAmount FROM Claim_B')
258 AS Claim_B(ClaimID INT, PolicyID INT, Type VARCHAR(50), Status VARCHAR(20), ClaimedAmount DECIMAL(10,2));
259
260 --Validate with COUNT and CHECKSUM
261 ---Count rows on Node_A (local)
262
263 SELECT COUNT(*) AS count_a FROM Claim_A;

```

Data Output Messages Notifications

Showing rows: 1 to 5

Page No:

| | claimid integer | policyid integer | type character varying (50) | status character varying (20) | claimedamount numeric (10,2) |
|---|--------------------|---------------------|--------------------------------|----------------------------------|---------------------------------|
| 1 | 1 | 4 | Health | Approved | 50000.00 |
| 2 | 2 | 5 | Home | Pending | 70000.00 |
| 3 | 3 | 6 | Life | Approved | 110000.00 |
| 4 | 4 | 7 | Auto | Closed | 45000.00 |
| 5 | 5 | 8 | Travel | Approved | 60000.00 |

A1.4

```

260 --Validate with COUNT and CHECKSUM
261 ---Count rows on Node_A (local)
262
263 SELECT COUNT(*) AS count_a FROM Claim_A;
264
265 ---Count rows on Node_B (remote)
266
267 SELECT COUNT(*) AS count_b
268 FROM dblink('proj_link', 'SELECT COUNT(*) FROM Cla
269
270 ---Count rows in the unified view
271
272 SELECT COUNT(*) AS total_count FROM Claim_ALL;
273
274
275

```

Data Output Messages Notifications

```

ERROR:  relation "claim_all" does not exist
LINE 1: SELECT COUNT(*) AS total_count FROM Claim_ALL;

```

```

SQL state: 42P01
Character: 37

```

A2 :Database Link & Cross-Node Join (3–10 rows result)

WHAT TO DO

1. From Node_A, create database link 'proj_link' to Node_B.
2. Run remote SELECT on Policy@proj_link showing up to 5 sample rows.
3. Run a distributed join: local Claim_A (or base Claim) joined with remote Agent@proj_link returning between 3 and 10 rows total; include selective predicates to stay within the row budget.

EXPECTED OUTPUT

- ✓ - CREATE DATABASE LINK proj_link with connection details.
- ✓ - Screenshot of SELECT * FROM Policy@proj_link FETCH FIRST 5 ROWS ONLY.
- ✓ - Screenshot of distributed join on Claim ⋈ Agent@proj_link returning 3–10 rows.

A2.1

```
- 2 | Create database  
- Replace host, dbnam  
- =====  
ELECT dblink_connect(  
    'proj link'.
```

A2.2.

```
308 -- Run remote SELECT on Policy@proj_link (up to 5 sample rows)
309 SELECT *
310 FROM dblink('proj_link',
311            'SELECT PolicyID, ClientID, AgentID, Type, Premium, StartDate, EndDate, Status FROM Policy LIMIT 5')
312 AS Policy_Remote(
313     PolicyID INT,
314     ClientID INT,
315     AgentID INT,
316     Type VARCHAR(50),
317     Premium DECIMAL(10,2),
318     StartDate DATE,
319     EndDate DATE,
320     Status VARCHAR(20)
321 );
322
323 -- =====
```

Data Output Messages Notifications

| Showing rows: 1 to 5 Page No: 1 | | | | | | | | | |
|----------------------------------|------------------|------------------|-----------------|-----------------------------|------------------------|----------------|--------------|-------------------------------|--|
| | policyid integer | clientid integer | agentid integer | type character varying (50) | premium numeric (10,2) | startdate date | enddate date | status character varying (20) | |
| 1 | 1 | 1 | 1 | Home | 300000.00 | 2025-03-01 | 2026-03-01 | Active | |
| 2 | 2 | 2 | 2 | Life | 500000.00 | 2025-04-01 | 2030-04-01 | Active | |
| 3 | 3 | 3 | 1 | Travel | 120000.00 | 2025-05-01 | 2026-05-01 | Active | |
| 4 | 4 | 1 | 1 | Home | 300000.00 | 2025-03-01 | 2026-03-01 | Active | |
| 5 | 5 | 2 | 2 | Life | 500000.00 | 2025-04-01 | 2030-04-01 | Active | |

A2.3

```
324 -- Run a distributed join: local Claim_A joined with remote Agent@proj_link
325 -- Includes selective predicates to keep row count between 3-10
326 -- =====
327 SELECT
328     c.ClaimID,
329     c.PolicyID,
330     c.Type AS ClaimType,
331     c.Status AS ClaimStatus,
332     c.ClaimedAmount,
333     a.FullName AS AgentName,
334     a.Contact AS AgentContact,
335     a.Branch AS AgentBranch
336 FROM Claim_A c
337 JOIN dblink(
338     'proj_link',
339     'SELECT AgentID, FullName, Contact, Branch, Experience FROM Agent WHERE Experience >= 2 LIMIT 5'
340 )
341 AS a(
342     AgentID INT,
343     FullName VARCHAR(100),
344     Contact VARCHAR(50),
345     Branch VARCHAR(50),
346     Experience INT
347 )
```

Data Output Messages Notifications

| Showing rows: 1 to 4 Page No: 1 of 1 | | | | | | | | |
|---------------------------------------|----------------------|------------------|----------------------------------|------------------------------------|------------------------------|-----------------------------------|-------------------------------------|-------------------------------------|
| | claimid [PK] integer | policyid integer | claimtype character varying (50) | claimstatus character varying (20) | claimedamount numeric (10,2) | agentname character varying (100) | agentcontact character varying (50) | agentbranch character varying (100) |
| 1 | 4 | 1 | Life | Approved | 120000.00 | Mary Uwase | 0789000002 | Musanze |
| 2 | 1 | 1 | Health | Approved | 80000.00 | Mary Uwase | 0789000002 | Musanze |
| 3 | 2 | 2 | Auto | Pending | 60000.00 | Eric Habimana | 0789000003 | Huye |
| 4 | 3 | 3 | Home | Closed | 90000.00 | Mary Uwase | 0789000002 | Musanze |

A3 :Parallel vs Serial Aggregation (≤ 10 rows data)

WHAT TO DO

1. Run a SERIAL aggregation on Claim_ALL over the small dataset (e.g., totals by a domain column). Ensure result has 3–10 groups/rows.
2. Run the same aggregation with `/*+ PARALLEL(Claim_A,8) PARALLEL(Claim_B,8) */` to force a parallel plan despite small size.
3. Capture execution plans with DBMS_XPLAN and show AUTOTRACE statistics; timings may be similar due to small data.
4. Produce a 2-row comparison table (serial vs parallel) with plan notes.

EXPECTED OUTPUT

- ✓ - Two SQL statements (serial and parallel) with hints.
- ✓ - DBMS_XPLAN outputs for both runs (showing parallel plan chosen in the hinted version).
- ✓ - AUTOTRACE / timing evidence and a small comparison table (mode, ms, buffer gets).

A3.1

```
373 -- 3 MANUAL COMPARISON TABLE
374 -- Fill in based on EXPLAIN ANALYZE output
375 SELECT 'Serial' AS Execution_Type, '5-10' AS Total_Rows,
376        'Aggregate node; single worker; no parallelism' AS Plan_Notes
377 UNION ALL
378 SELECT 'Parallel', '5-10',
379        'Gather + Parallel Seq Scan + Parallel Aggregate; may still be serial on small dataset';
380
```

Data Output Messages Notifications

Showing rows: 1 to 2

| | execution_type text | total_rows text | plan_notes text |
|---|------------------------|--------------------|---|
| 1 | Serial | 5-10 | Aggregate node; single worker; no parallelism |
| 2 | Parallel | 5-10 | Gather + Parallel Seq Scan + Parallel Aggregate; may still be serial on small dataset |

A3.2

```
399 -- 2 PARALLEL AGGREGATION ON Claim_ALL
400 -- Enable parallel execution and allow up to 8 parallel workers
401 ALTER SESSION ENABLE PARALLEL DML;
402 ALTER SESSION FORCE PARALLEL_QUERY PARALLEL 8;
403
404 -- Run parallel aggregation and capture cursor
405 VAR parallel_cursor NUMBER;
406 BEGIN
407     OPEN :parallel_cursor FOR
408     SELECT /*+ PARALLEL(Claim_A,8) PARALLEL(Claim_B,8) */
409           Type, SUM(ClaimedAmount) AS total_claim
410     FROM Claim_ALL
411     GROUP BY Type
412     ORDER BY Type;
413 END;
414 /
415
416 -- Display execution plan and stats for parallel
417 SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY_CURSOR(FORMAT => 'ALLSTATS LAST'));
418
```

A3.3

```
415
416 -- Display execution plan and stats for parallel
417 SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY_CURSOR(FORMAT => 'ALLSTATS LAST'));
418 |
```

A3.4

```
419 --- PRODUCE COMPARISON TABLE (SERIAL vs PARALLEL)
420 |
421 -- Example: manually create comparison notes
422 -- In practice, you may populate this from your DBMS_XPLAN output
423 SELECT 'Serial' AS Execution_Type, '5-10' AS Total_Rows,
424        'HashAggregate; single worker; no parallelism' AS Plan_Notes
425 FROM dual
426 UNION ALL
427 SELECT 'Parallel' AS Execution_Type, '5-10' AS Total_Rows,
428        'Parallel HashAggregate; Gather node; may still execute serially on small dataset' AS Plan_Notes
429 FROM dual;
430
431
```


A4 :Two-Phase Commit & Recovery (2 rows)

WHAT TO DO

1. Write one PL/SQL block that inserts ONE local row (related to Claim) on Node_A and ONE remote row into Payment@proj_link (or Claim@proj_link); then COMMIT.
2. Induce a failure in a second run (e.g., disable the link between inserts) to create an in-doubt transaction; ensure any extra test rows are ROLLED BACK to keep within the ≤ 10 committed row budget.
3. Query DBA_2PC_PENDING; then issue COMMIT FORCE or ROLLBACK FORCE; re-verify consistency on both nodes.
4. Repeat a clean run to show there are no pending transactions.

EXPECTED OUTPUT

- ✓ - PL/SQL block source code (two-row 2PC).
- ✓ - DBA_2PC_PENDING snapshot before/after FORCE action.
- ✓ - Final consistency check: the intended single row per side exists exactly once; total committed rows remain ≤ 10 .

A4.1

```
39  -- 1 Ensure local table exists
40
41  CREATE TABLE IF NOT EXISTS Claim_A (
42      ClaimID SERIAL PRIMARY KEY,
43      PolicyID INT NOT NULL,
44      Type VARCHAR(50),
45      Status VARCHAR(20),
46      ClaimedAmount DECIMAL(10,2)
47  );
48
```










ata Output Messages Notifications

OTICE: relation "claim_a" already exists, skipping
REATE TABLE

A4.2

```
520 -- CHECK PENDING 2PC TRANSACTIONS
521
522 SELECT * FROM pg_prepared_xacts;
```

Data Output Messages Notifications

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|------------------|
|  |  |  |  |  |  |  |  |  | SQL |
| | transaction xid | | gid text | | prepared timestamp with time zone | | owner name | | database name |

A4.3

```
-- Force commit or rollback
COMMIT PREPARED 'tx_local...';
-- or
ROLLBACK PREPARED 'tx_local...';

-- Use dblink to COMMIT/ROLLBACK prepared
SELECT dblink_exec('proj_link', 'COMMIT PREPARED 'tx_remote...');');
```

A4.4

```
559 -- Remote check via dblink
560 SELECT * FROM dblink('proj_link', 'SELECT * FROM Payment WHERE PolicyID=4')
561 AS t(PaymentID INT, PolicyID INT, Amount DECIMAL);
562
563
564
565
566
```

Data Output Messages Notifications

ERROR: could not establish connection
missing "=" after "proj_link" in connection info string

A5 :Distributed Lock Conflict & Diagnosis (no extra rows)

WHAT TO DO

1. Open Session 1 on Node_A: UPDATE a single row in Claim or Payment and keep the transaction open.
2. Open Session 2 from Node_B via Claim@proj_link or Payment@proj_link to UPDATE the same logical row.
3. Query lock views (DBA_BLOCKERS/DBA_WAITERS/V\$LOCK) from Node_A to show the waiting session.
4. Release the lock; show Session 2 completes. Do not insert more rows; reuse the existing ≤ 10 .

EXPECTED OUTPUT

- ✓ - Two UPDATE statements showing the contested row keys.
- ✓ - Lock diagnostics output identifying blocker/waiter sessions.
- ✓ - Timestamps showing Session 2 proceeds only after lock release.

A5.1

```
564
565 -- On Node_A (run in a psql session)
566 CREATE TABLE IF NOT EXISTS Claim_A (
567     ClaimID SERIAL PRIMARY KEY,
568     PolicyID INT NOT NULL,
569     Type VARCHAR(50),
570     Status VARCHAR(20),
571     ClaimedAmount DECIMAL(10,2)
572 );
573
574 -- Ensure a row exists to reuse (ClaimID = 1)
575 INSERT INTO Claim_A (PolicyID, Type, Status, ClaimedAmount
576 SELECT 1, 'Health', 'Approved', 80000
577 WHERE NOT EXISTS (SELECT 1 FROM Claim_A WHERE ClaimID = 1)
578
579
580
```

Data Output [Messages](#) Notifications

INSERT 0 0

Query returned successfully in 6 secs 205 msec.

A5.2

```
564
565 -- On Node_A (run in a psql session)
566 CREATE TABLE IF NOT EXISTS Claim_A (
567     ClaimID SERIAL PRIMARY KEY,
568     PolicyID INT NOT NULL,
569     Type VARCHAR(50),
570     Status VARCHAR(20),
571     ClaimedAmount DECIMAL(10,2)
572 );
573
574 -- Ensure a row exists to reuse (ClaimID = 1)
575 INSERT INTO Claim_A (PolicyID, Type, Status, ClaimedAmount
576 SELECT 1, 'Health', 'Approved', 80000
577 WHERE NOT EXISTS (SELECT 1 FROM Claim_A WHERE ClaimID = 1)
578
579
580
```

Data Output Messages Notifications

INSERT 0 0

Query returned successfully in 6 secs 205 msec.

A5.3a

```
598 -- 1) View current locks and who is waiting
599 -- Shows pg_locks joined to pg_stat_activity
600 SELECT
601     a.pid AS pid,
602     a.username,
603     a.state,
604     a.query AS active_query,
605     l.locktype,
606     l.mode,
607     l.granted
608 FROM pg_locks l
609 JOIN pg_stat_activity a ON l.pid = a.pid
610 WHERE (l.locktype IN ('tuple','relation') OR l.mode LIKE '%Exclusive%')
611 ORDER BY a.pid;
```

Data Output Messages Notifications

| | pid integer | username name | state text | active_query text |
|---|----------------|------------------|---------------|--|
| 1 | 31031 | postgres | active | -- 1) View current locks and who is waiting -- Shows pg_locks joined to pg_stat_activity SELECT a.pid AS pid, a.username, a.state, a.query AS active_query, l.locktype, l.m... |
| 2 | 31031 | postgres | active | -- 1) View current locks and who is waiting -- Shows pg_locks joined to pg_stat_activity SELECT a.pid AS pid, a.username, a.state, a.query AS active_query, l.locktype, l.m... |
| 3 | 31031 | postgres | active | -- 1) View current locks and who is waiting -- Shows pg_locks joined to pg_stat_activity SELECT a.pid AS pid, a.username, a.state, a.query AS active_query, l.locktype, l.m... |
| 4 | 31031 | postgres | active | -- 1) View current locks and who is waiting -- Shows pg_locks joined to pg_stat_activity SELECT a.pid AS pid, a.username, a.state, a.query AS active_query, l.locktype, l.m... |
| 5 | 31031 | postgres | active | -- 1) View current locks and who is waiting -- Shows pg_locks joined to pg_stat_activity SELECT a.pid AS pid, a.username, a.state, a.query AS active_query, l.locktype, l.m... |
| 6 | 31031 | postgres | active | -- 1) View current locks and who is waiting -- Shows pg_locks joined to pg_stat_activity SELECT a.pid AS pid, a.username, a.state, a.query AS active_query, l.locktype, l.m... |
| 7 | 31031 | postgres | active | -- 1) View current locks and who is waiting -- Shows pg_locks joined to pg_stat_activity SELECT a.pid AS pid, a.username, a.state, a.query AS active_query, l.locktype, l.m... |

A5.3b

```
613 -- 2) Show waiters vs blockers (blocking_pid is in pg_stat_activity.waiting? In PG12+ use wait_event_type/wait_event)
614 -- A typical approach: see which backends are waiting and which are blocking
615 SELECT
616     blocked.pid AS blocked_pid,
617     blocked.query AS blocked_query,
618     blocking.pid AS blocking_pid,
619     blocking.query AS blocking_query
620 FROM pg_stat_activity blocked
621 JOIN pg_locks bl ON blocked.pid = bl.pid AND NOT bl.granted
622 JOIN pg_locks kl ON (bl.locktype = kl.locktype AND bl.database IS NOT DISTINCT FROM kl.database AND bl.relation IS NOT DISTINCT FROM kl.r
623 JOIN pg_stat_activity blocking ON kl.pid = blocking.pid AND kl.granted
624 WHERE blocked.pid <> blocking.pid
625     AND blocked.query IS NOT NULL
626     AND blocking.query IS NOT NULL;
627
628
629
```

Data Output Messages Notifications

| blocked_pid integer | blocked_query text | blocking_pid integer | blocking_query text |
|------------------------|-----------------------|-------------------------|------------------------|
|------------------------|-----------------------|-------------------------|------------------------|